Home exam 1 - IN3200

Candidate nr: 15129
(Dated: March 16, 2020)

I. INTRODUCTION

In this report we look into the main algorithmic aspects of the code implementations and present time measurements of the serial and parallelized codes.

II. BACKGROUND

Webgraphs

Webgraphs describe hyperlinks between to webpages i and j. Each webpage is represented by a node and the hyperlink between them is known as an edge. In this report I'll define N as the number of nodes and N_{links} as the number of edges in a *directed* webgraph. By directed, we mean that the hyperlink $j \to i$ is distinct from $j \to i$.

Hyperlink matrix representation of webgraphs

Webgraphs can be represented by what is called a hyperlink matrix A. An element $A_{ij} = 1$ if there exists a hyperlink $i \to j$. If no such link exists, $A_{ij} = 0$.

Compressed row storage (CRS) of the hyperlink matrix

For large webgraphs, it's expected that the hyperlink matrix will consist mainly of zeros. Since the storage in the matrix format scales with N^2 where N is the number of nodes, storing the data becomes a bottleneck. The CRS format is based on to different arrays, a row pointer storing how many row elements that accumulates of a set of rows n and column index array that stores the column index for a given element. Let r denote the row pointer. The general form of the row pointer is

$$r = (0, r_0, r_0 + r_1, ..., r_0 + r_1 + \cdots + r_{N-1}), \quad (1)$$

where r_i is the number of row elements of row i. The row pointer thus has length $N + 1$. The column index array is of length N_{links} and stores the column indices for each row. To extract how many row elements there are row i, we take the difference $r_i - r_{i-1}$.

Mutual web linkages and number of involvements

Mutual web linkages is defined as the number of times to outbound nodes i and j are directly linked to a inbound node k, with $i \neq j \neq k$. That is, it's a count of how many times the exists a hyperlink $i \to k$ and $j \to k$ simultaneously. The *number of involvements* a given node has is defined as the number of times a given node is involved as an outbound node. That is how many times a node i is involved in $i \to k$ with $j \to k$.

III. IMPLEMENTATION

Reading the Webgraph from file

1. Hyperlink matrix storage

To extract the data to store in a hyperlink matrix is straight forward. The following code snippet shows how I did it.

```
int FromNodeId, ToNodeId;
for (int k = 0; k < N_links; k++){
  fscanf(fp, "%d %d", &FromNodeId, &ToNodeId);
  (*table2D)[ToNodeId][FromNodeId] = (char) 1;
}
```

2. CRS storage

To extract the data from the file into CRS storage, we first must sort the arrays. I chose to sort this using the *shellsort* algorithm with gap $= N/2$. The following code snippet shows the implementation.

```
int tmp1, tmp2, i, j, gap;
for (gap = *N_links/2; gap > 0; gap /= 2){
  for (i = gap; i < *N_links; i++){
  tmp1 = row_elems[i];
  tmp2 = (*col_idx)[i];
  for (j = i; j >= gap && row_elems[j-gap] >
      tmp1; j -= gap){
    row_elems[j] = row_elems[j-gap];
    (*col_idx)[j] = (*col_idx)[j-gap];
  }
  row_elems[j] = tmp1;
  (*col_idx)[j] = tmp2;
  }
}
```

Here, row_elems is just there to temporarily store row node ids. The row pointer is simply made using an array counting how many elements each row has consistent with eq. (1). The following code demonstrates this.

```
*row_ptr = (int*)calloc(*N+1, sizeof(int*));
int count = 0;
for (int i = 0; i < *N; i++){
  count += row_count[i];
  (*row_ptr)[i+1] = count;
}
```

Counting mutual web links

3. Counting mutual web links with the hyperlink matrix

To count the number of web links that each node is an outbound participant is in principle a simple matter when we use the hyperlink matrix representation. It is simply to check when $A_{ij} = A_{ik} = 1$ with $i \neq j \neq k$. However, to make the code implementation efficient, it's important to do this economically. For a given node i, we should check every $j \neq i$ and only if $A_{ij} = 1$ should we check whether $A_{ik} = 1$. The following code demonstrates how this can be implemented. It also counts the total number of web linkages which is just the total number of times mutual web links occur.

```
for (i = 0; i < N; i++){
  for (j = 0; j < N; j++){
    if (table2D[i][j] == 1){
      for (k = j + 1; k < N; k++){
      if (table2D[i][k] == 1)
        {
          total_mutual_web_linkages++;
          num_involvements[j]++;
          num_involvements[k]++;
        }
      }
    }
  }
}
```

4. Counting mutual web links with the CRS format

To count the total number of web links in this format, we can reduce the number of floating point operations by computing the sum analytically. To find the contribution a given row i has to the total number of web linkages, let's first define m_i to be the number of elements on row i such that $m_i = r_i - r_{i-1}$. Then

$$S = \sum_{k=0}^{m_i-1} k = \frac{m_i(m_i - 1)}{2}, \qquad (2)$$

where the standard formula for the sum of the $m_i - 1$ first integers were used. The following code snippet shows the implementation.

```
for (int i = 0; i < N; i++){
  tmp = row_ptr[i];
  row_elems = row_ptr[i+1]-tmp;
  for (int j = 0; j < row_elems; j++){
    num_involvements[col_idx[j+tmp]] +=
        row_elems-1;
  }
  total_mutual_web_linkages += (row_elems)*(
      row_elems-1);
}
total_mutual_web_linkages *= 0.5;
```

As can be seen from the code, a tmp variable is defined to avoid loading row_ptr[i] in j-dependent loop. In addition we use the sum formula instead of computing the actual sum for each i manually.

Finding top webpages

First of all, by top webpages we mean the webpages or nodes that are most involved as an outbound node in mutual linkages. To find the n top webpages in a webgraph, I've used shellsort again, but this time sorting in descending order. For small n this isn't the wisest choice since the algorithm is only dependent on how many nodes there are. However, for larger n's it's an efficient choice. The code implementation is as follows.

```
int gap, i, j, tmp1, tmp2;
  for (gap = num_webpages/2; gap > 0; gap /=
      2){
  for (i = gap; i < num_webpages; i++){
    tmp1 = num_involvements[i];
    tmp2 = webpage_number[i];
    for (j = i; j >= gap && num_involvements[j
        -gap] < tmp1; j -= gap){
      num_involvements[j] = num_involvements[j
          -gap];
      webpage_number[j] = webpage_number[j-gap
          ];
    }
    num_involvements[j] = tmp1;
    webpage_number[j] = tmp2;
  }
}
```

Parallelization of count_mutual_links1

The function count_mutual_links1 can easily be parallelized by a simple #pragma omp parallel for directive. However, all threads need access to the entire num_involvements array which can give rise to race conditions. We must therefore also insert a #pragma omp critical region inside the inner loop. The following code demonstrates this.

```
int i, j, k;
#pragma omp parallel for private(i,j,k)
    reduction(+:total_mutual_web_linkages)
for (i = 0; i < N; i++){
  for (j = 0; j < N; j++){
    if (table2D[i][j] == 1){
      for (k = j + 1; k < N; k++){
        if (table2D[i][k] == 1)
          {
            total_mutual_web_linkages++;
            #pragma omp critical
            {
              num_involvements[j]++;
              num_involvements[k]++;
            }
          }
      }
    }
  }
}
```

Parallelization of count_mutual_links2

This function is also readily parallelizable, but also here we must avoid race conditions. Here, however, we can use the #pragma omp atomic directive instead. The following code shows the parallelized version.

```
int tmp, row_elems;
#pragma omp parallel for private(tmp,
    row_elems) reduction(+:
    total_mutual_web_linkages)
for (int i = 0; i < N; i++){
  tmp = row_ptr[i];
  row_elems = row_ptr[i+1] - tmp;
  for (int j = 0; j < row_elems; j++){
    //Insert atomic to avoid race conditions
        when updating num_involvements.
    #pragma omp atomic
    num_involvements[col_idx[j+tmp]] +=
        row_elems-1;
  }
  total_mutual_web_linkages += (row_elems)*(
      row_elems-1);
}
```

IV.   RESULTS

Timing of serial codes

The measured time of the serial implementations of the various functions are shown in table I.

Function name	Time in seconds
read_graph_from_file1	0.083591
count_mutual_links1	0.885690
read_graph_from_file2	0.356536
count_mutual_links2	0.001885
top_n_webpages	0.019348

TABLE I. The table shows the measured time using clock() from the Ctime-library. read_graph_from_file1 and count_mutual_links1 was applied to a web-graph containing $N = 10000$ nodes and $N_{links} = 37841$ edges as found in the file test_webpages.txt. The data in this file was extracted from web-NotreDame.txt. read_graph_from_file2 and count_mutual_links2 was applied directly to the web-graph contained in web-NotreDame.txt. This file contained $N = 325729$ nodes and $N_{links} = 1479143$ edges.

Parallelized version of count_mutual_links1

Using OpenMP to parallelize count_mutual_links1 gave the time measurements shown in figure 1.

FIG. 1.   The figure shows the time used in seconds by count_mutual_links1 on a webgraph consisting of $N = 50000$ nodes and $N_{links} = 146823$ edges. The webgraph was extracted from web-NotreDame.txt. The red points show the actual measured datapoints.

The speedup relative to one thread is shown in figure 2.

FIG. 2. The figure shows speedup relative to a single thread by count_mutual_links1 on a webgraph consisting of $N = 50000$ nodes and $N_{links} = 146823$ edges. The webgraph was extracted from web-NotreDame.txt. The red points show the actual measured datapoints.

Parallelized version of count_mutual_links2

Using OpenMP to parallelize count_mutual_links2 and measuring the time used by the function for different number of threads yielded the results shown in figure 3.

FIG. 3. The figure shows the time used in seconds by count_mutual_links2 on a webgraph consisting of $N = 325729$ nodes and $N_{\text{links}} = 1479143$ edges. The webgraph is found in the file web-NotreDame.txt

The speedup relative to a single thread is shown in figure 4.

FIG. 4. The figure shows speedup relative to a single thread by count_mutual_links1 on a webgraph consisting of $N = 325729$ nodes and $N_{\text{links}} = 1479143$ edges. The webgraph is from the file web-NotreDame.txt. The red points show the actual measured datapoints.

V.   CONCLUDING REMARKS

As we can see from figure 1, increasing the number of threads acquires a significant speed up by increasing the number of threads. In the matrix representation, it makes sense to use threads in the range of $[500, 1000]$. When using the CRS format however, the sweetspot is not so easily deduced, but figure 3 indicates that choosing the thread number to be about 8 is the optimal solution.